

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 6.1: Mining Data Streams

Data Streams

- ❖ In many data mining situations, we do not know the entire data set in advance
- ❖ Stream Management is important when the input rate is controlled **externally**:
 - Google queries
 - Twitter or Facebook status updates
- ❖ We can think of the **data** as **infinite** and **non-stationary** (the distribution changes over time)

Characteristics of Data Streams

- ❖ Traditional DBMS: data stored in *finite, persistent data sets*
- ❖ Data Streams: distributed, continuous, unbounded, rapid, time varying, noisy, . . .
- ❖ Characteristics
 - Huge volumes of continuous data, possibly infinite
 - Fast changing and requires fast, real-time response
 - Random access is expensive—single scan algorithm (can only have one look)
 - Store only the summary of the data seen thus far

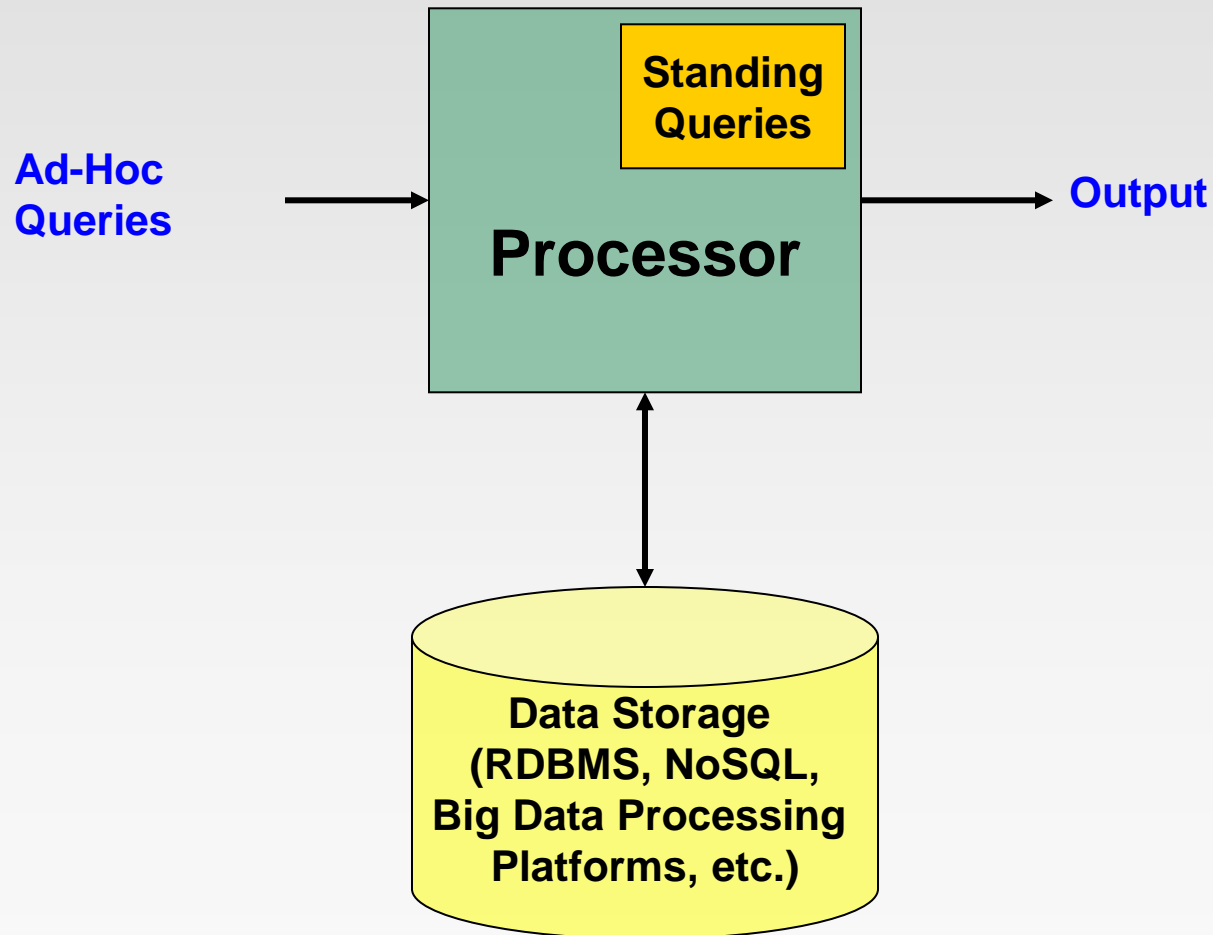
Massive Data Streams

- ❖ Data is *continuously growing* faster than our ability to store or index it
- ❖ There are 3 Billion Telephone Calls in US each day, 30 Billion emails daily, 1 Billion SMS, IMs
- ❖ Scientific data: NASA's observation satellites generate billions of readings each per day
- ❖ IP Network Traffic: up to 1 Billion packets per hour per router. Each ISP has many (hundreds) routers!
- ❖

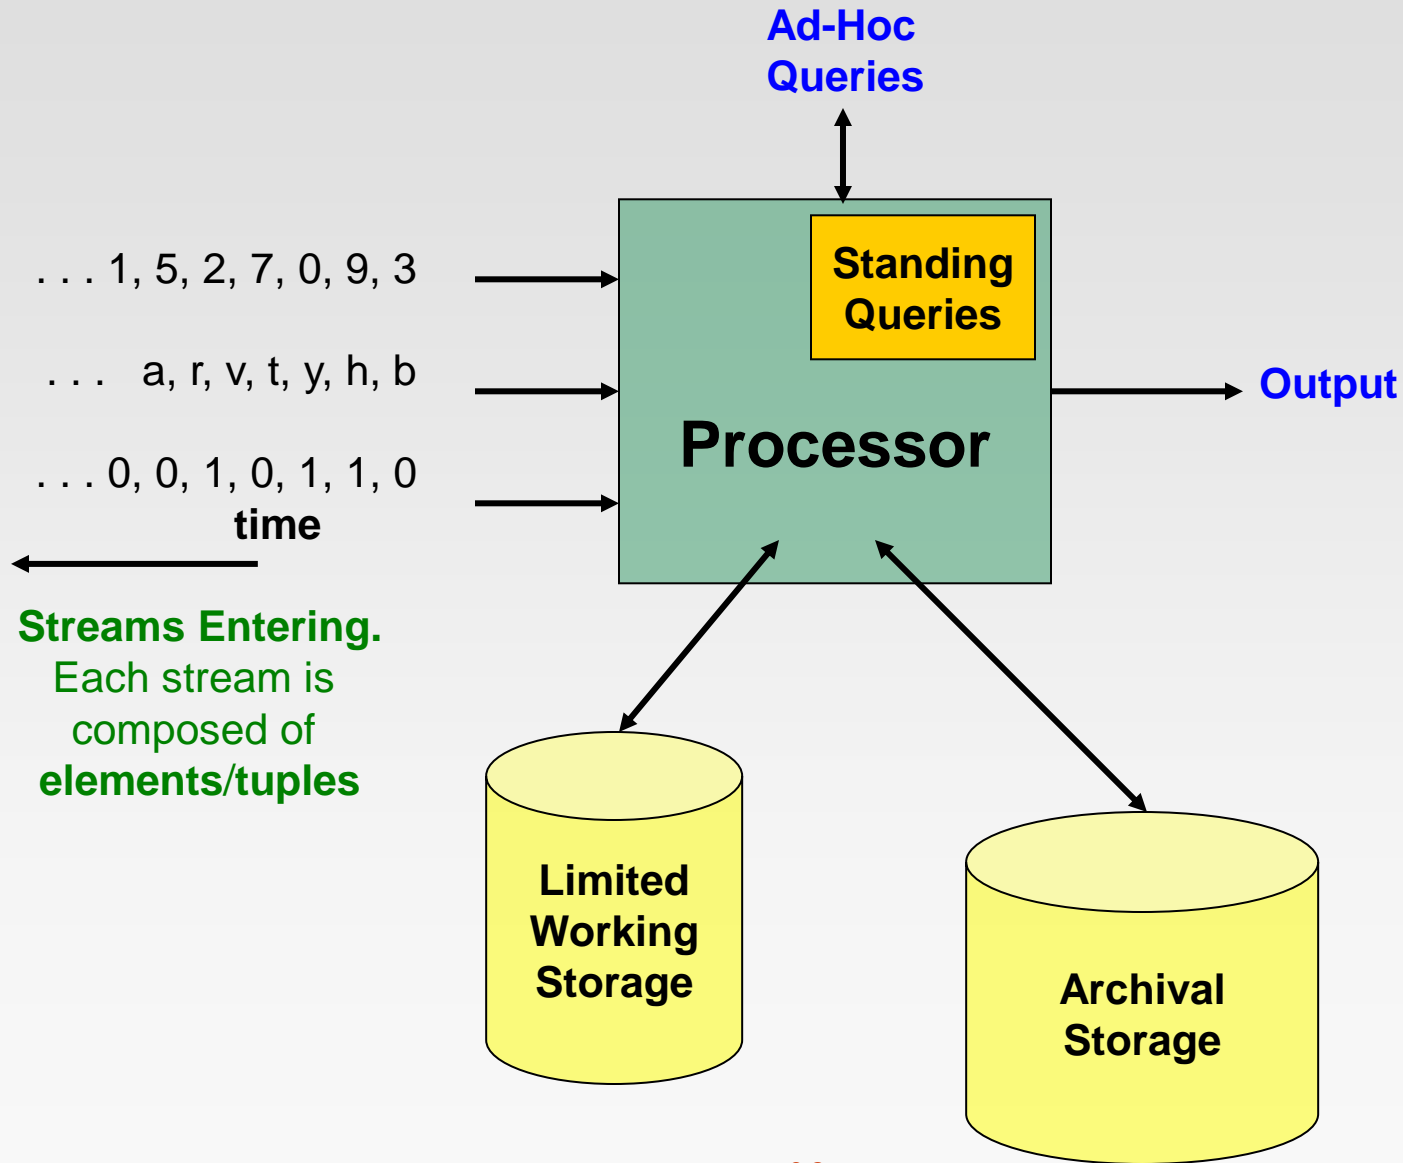
The Stream Model

- ❖ Input elements enter at a rapid rate, at one or more input ports (i.e., streams)
 - We call elements of the stream tuples
- ❖ The system cannot store the entire stream accessibly
- ❖ **Q: How do you make critical calculations about the stream using a limited amount of memory?**

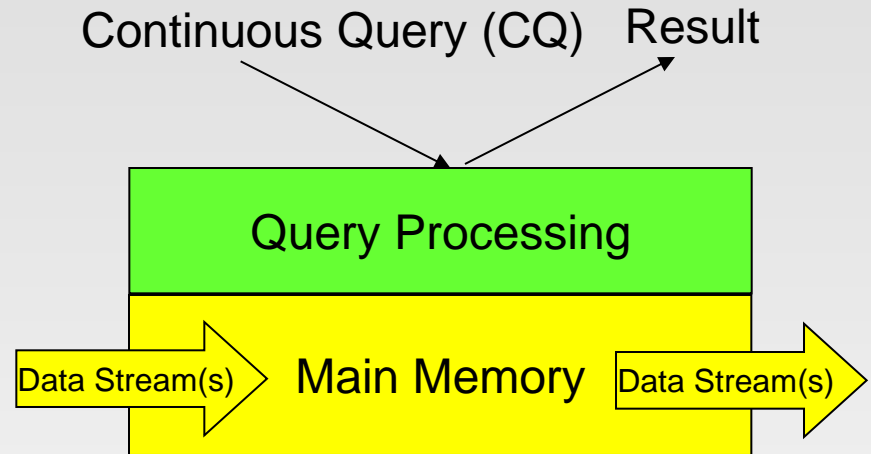

Database Management System (DBMS) Data Processing



General Data Stream Management System (DSMS) Processing Model



DBMS vs. DSMS #1



DBMS vs. DSMS #2

❖ Traditional DBMS:

- stored sets of relatively **static** records with no pre-defined notion of time
- good for applications that require **persistent data storage** and **complex querying**

❖ DSMS:

- support on-line analysis of rapidly **changing** data streams
- data stream: real-time, continuous, ordered (implicitly by arrival time or explicitly by timestamp) sequence of items, too large to store entirely, no ending
- **continuous** queries

DBMS vs. DSMS #3

DBMS

- ❖ Persistent relations
(relatively static, stored)
- ❖ One-time queries
- ❖ Random access
- ❖ “Unbounded” disk store
- ❖ Only current state matters
- ❖ No real-time services
- ❖ Relatively low update rate
- ❖ Data at any granularity
- ❖ Assume precise data
- ❖ Access plan determined by query processor, physical DB design

DSMS

- ❖ Transient streams
(on-line analysis)
- ❖ Continuous queries (CQs)
- ❖ Sequential access
- ❖ Bounded main memory
- ❖ Historical data is important
- ❖ Real-time requirements
- ❖ Possibly multi-GB arrival rate
- ❖ Data at fine granularity
- ❖ Data stale/imprecise
- ❖ Unpredictable/variable data arrival and characteristics


Problems on Data Streams

- ❖ Types of queries one wants on answer on a data stream: (we'll learn these today)
 - Sampling data from a stream
 - ▶ Construct a random sample
 - Queries over sliding windows
 - ▶ Number of items of type x in the last k elements of the stream
 - Filtering a data stream
 - ▶ Select elements with property x from the stream
 - Counting distinct elements
 - ▶ Number of distinct elements in the last k elements of the stream
 - Finding frequent elements
 -

Applications

- ❖ Mining query streams
 - Google wants to know what queries are more frequent today than yesterday
- ❖ Mining click streams
 - Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour
- ❖ Mining social network news feeds
 - E.g., look for trending topics on Twitter, Facebook
- ❖ Sensor Networks
 - Many sensors feeding into a central controller
- ❖ Telephone call records
 - Data feeds into customer bills as well as settlements between telephone companies
- ❖ IP packets monitored at a switch
 - Gather information for optimal routing

Example: IP Network Data




- ❖ Networks are sources of massive data: the **metadata** per hour per IP router is gigabytes
- ❖ Fundamental problem of data stream analysis:
 - Too much information to store or transmit
- ❖ So process data as it arrives
 - One pass, small space: the data stream approach
- ❖ **Approximate answers** to many questions are OK, if there are guarantees of result quality

Part 1: Sampling Data Streams

Sampling from a Data Stream

- ❖ Since we can not store the entire stream, one obvious approach is to store a **sample**



- ❖ Two different problems:
 - (1) Sample a **fixed proportion** of elements in the stream (say 1 in 10)
 - ▶ As the stream grows the sample also gets bigger
 - (2) Maintain a **random sample of fixed size** over a potentially infinite stream
 - ▶ As the stream grows, the sample is of fixed size
 - ▶ At any “time” t we would like a random sample of s elements
 - What is the property of the sample we want to maintain?
For all time steps t , each of t elements seen so far has equal probability of being sampled

Sampling a Fixed Proportion

- ❖ Problem 1: Sampling fixed proportion
- ❖ Scenario: Search engine query stream
 - **Stream of tuples:** (user, query, time)
 - **Answer questions such as:** How often did a user run the same query in a single days
 - Have space to store $1/10^{\text{th}}$ of query stream
- ❖ **Naïve solution:**
 - Generate a random integer in **[0..9]** for each query
 - Store the query if the integer is **0**, otherwise discard

Problem with Naïve Approach

❖ **Simple question:** What fraction of queries by an average search engine user are duplicates?

➤ Suppose each user issues x queries once and d queries twice (total of $x+2d$ queries)

➤ **Correct answer:** $d/(x+d)$

➤ **Proposed solution:** We keep 10% of the queries

➤ Sample will contain $x/10$ of the singleton queries and $2d/10$ of the duplicate queries at least once

➤ But only $d/100$ pairs of duplicates

— $d/100 = 1/10 \cdot 1/10 \cdot d$

➤ Of d “duplicates” $18d/100$ appear exactly once

— $18d/100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$

➤ **So the sample-based answer is** $\frac{\frac{d}{100}}{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}} = \frac{d}{10x+19d} \neq d/(x+d)$

Solution: Sample Users

Solution:

- ❖ Pick $1/10^{\text{th}}$ of **users** and take all their searches in the sample
- ❖ Use a hash function that hashes the username or user id uniformly into 10 buckets
 - We hash each username to one of ten buckets, 0 through 9
 - If the user hashes to bucket 0, then accept this search query for the sample, and if not, then not.

Generalized Problem and Solution

- ❖ Problem: Give a data stream, take a sample of fraction a/b .
- ❖ Stream of tuples with keys:
 - Key is some subset of each tuple's components
 - ▶ e.g., tuple is (user, search, time); key is **user**
 - Choice of key depends on application
- ❖ To get a sample of a/b fraction of the stream:
 - Hash each tuple's key uniformly into **b** buckets
 - Pick the tuple if its hash value is at most **a**



How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first 3 buckets

Sample Operator in Spark

- ❖ `sample(withReplacement, fraction, seed)`
 - Return a sampled subset of this RDD.
 - `withReplacement`: can elements be sampled multiple times
 - `fraction`: expected size of the sample as a fraction of this RDD's size without replacement
 - ▶ This is not guaranteed to provide exactly the fraction specified of the total count of the given
 - `seed`: seed for the random number generator

```
scala> val rdd = sc.parallelize(1 to 100)
rdd: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[90] at parallelize at <console>:31

scala> var sample1 = rdd.sample(true, 0.4, 2).collect
sample1: Array[Int] = Array(5, 5, 15, 19, 26, 27, 29, 38, 40, 45, 48, 48, 49, 50, 52, 54, 57, 58, 58, 59, 61, 67, 68, 68, 71, 73, 78, 82, 83, 83, 85, 88, 89, 89, 92, 95, 99)

scala> sample1.size
res7: Int = 37

scala> var sample2 = rdd.sample(false, 0.4, 2).collect
sample2: Array[Int] = Array(4, 5, 6, 7, 15, 21, 23, 26, 27, 36, 41, 42, 43, 44, 49, 50, 51, 52, 54, 61, 62, 66, 68, 77, 82, 86, 91, 93, 96)

scala> sample2.size
res8: Int = 29
```

Maintaining a Fixed-size Sample

- ❖ Problem 2: Fixed-size sample
- ❖ Suppose we need to maintain a random sample S of size exactly s tuples
 - E.g., main memory size constraint
- ❖ **Why?** Don't know length of stream in advance
- ❖ Suppose at time n we have seen n items
 - Each item is in the sample S with equal prob. s/n

How to think about the problem: say $s = 2$

Stream: a x c y z k c d e g...

At $n = 5$, each of the first 5 tuples is included in the sample S with equal prob.

At $n = 7$, each of the first 7 tuples is included in the sample S with equal prob.

Note that the same item is treated as different tuples at different timestamps

Impractical solution would be to store all the n tuples seen so far and out of them pick s at random

Solution: Fixed Size Sample

❖ Algorithm (a.k.a. Reservoir Sampling)

- Store all the first s elements of the stream to S
- Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - ▶ With probability s/n , keep the n^{th} element, else discard it
 - ▶ If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random

❖ Claim: This algorithm maintains a sample S with the desired property:

- After n elements, the sample contains each element seen so far with probability s/n

Proof: By Induction

❖ We prove this by induction:

- Assume that after n elements, the sample contains each element seen so far with probability s/n
- We need to show that after seeing element $n+1$ the sample maintains the property
 - ▶ Sample contains each element seen so far with probability $s/(n+1)$

❖ Base case:

- After we see $n=s$ elements the sample S has the desired property
 - ▶ Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction

- ❖ Inductive hypothesis: After n elements, the sample \mathbf{S} contains each element seen so far with prob. s/n
- ❖ **Now element $n+1$ arrives**
- ❖ **Inductive step:** For elements already in \mathbf{S} , probability that the algorithm keeps it in \mathbf{S} is:

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- ❖ So, at time n , tuples in \mathbf{S} were there with prob. s/n
- ❖ Time $n \rightarrow n+1$, tuple stayed in \mathbf{S} with prob. $n/(n+1)$
- ❖ So prob. tuple is in \mathbf{S} at time $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

takeSample Operator in Spark

- ❖ `takeSample(withReplacement, num, seed=None)`
 - Return a fixed-size sampled subset of this RDD.
 - `withReplacement`: can elements be sampled multiple times
 - `num`: sample size
 - This method should only be used if the resulting array is expected to be small, as all the data is loaded into the driver's memory.

```
scala> val rdd = sc.parallelize(1 to 100)
rdd: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[95] at parallel
ize at <console>:31

scala> var sample1 = rdd.takeSample(true, 20, 1)
sample1: Array[Int] = Array(67, 72, 29, 2, 37, 86, 16, 42, 68, 100, 46, 4,
  83, 67, 51, 69, 92, 24, 97, 8)

scala> sample1.size
res10: Int = 20
```

Part 2: Querying Data Streams

Sliding Windows

- ❖ A useful model of stream processing is that queries are about a **window** of length N – the N most recent elements received
- ❖ Interesting case: N is so large that the data cannot be stored in memory, or even on disk
 - Or, there are so many streams that windows for all cannot be stored
- ❖ Amazon example:
 - For every product X we keep 0/1 stream of whether that product was sold in the n -th transaction
 - We want answer queries, how many times have we sold X in the last k sales

Sliding Window: 1 Stream

- ❖ Sliding window on a single stream:

N = 7

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past Future →


Counting Bits (1)

❖ Problem:

- Given a stream of **0s** and **1s**
- Be prepared to answer queries of the form:
How many 1s are in the last k bits? where $k \leq N$

❖ Obvious solution:

- Store the most recent **N** bits
 - ▶ When new bit comes in, discard the **$N+1^{\text{st}}$** bit



Suppose $N=7$


Counting Bits (2)

- ❖ You can not get an exact answer without storing the entire window
- ❖ Real Problem:
What if we cannot afford to store N bits?
 - **E.g.**, we're processing 1 billion streams and **$N = 1$ billion**
- ❖ But we are happy with an approximate answer



An attempt: Simple solution

- ❖ Q: How many 1s are in the last N bits?
- ❖ A simple solution that does not really solve our problem: **Uniformity Assumption**



- ❖ Maintain 2 counters:
 - **S**: number of 1s from the beginning of the stream
 - **Z**: number of 0s from the beginning of the stream
- ❖ How many 1s are in the last N bits? $N \cdot \frac{S}{S+Z}$
- ❖ **But, what if stream is non-uniform?**
 - What if distribution changes over time?


The Datar-Gionis-Indyk-Motwani (DGIM) Algorithm

- ❖ Maintaining Stream Statistics over Sliding Windows (SODA'02)
- ❖ DGIM solution that does not assume uniformity
- ❖ We store $O(\log^2 N)$ bits per stream
 - If $N = 2^{16}$ (64KB), $\log(\log N) = \log(16) = 4$
- ❖ Solution gives approximate answer, never off by more than 50%
 - Error factor can be reduced to any fraction > 0 , with more complicated algorithm and proportionally more stored bits

Idea: Exponential Windows

- ❖ Solution that doesn't (quite) work:
 - Summarize **exponentially increasing** regions of the stream, looking backward
 - Drop small regions if they begin at the same point as a larger region

Window of width 16 has 6 1s



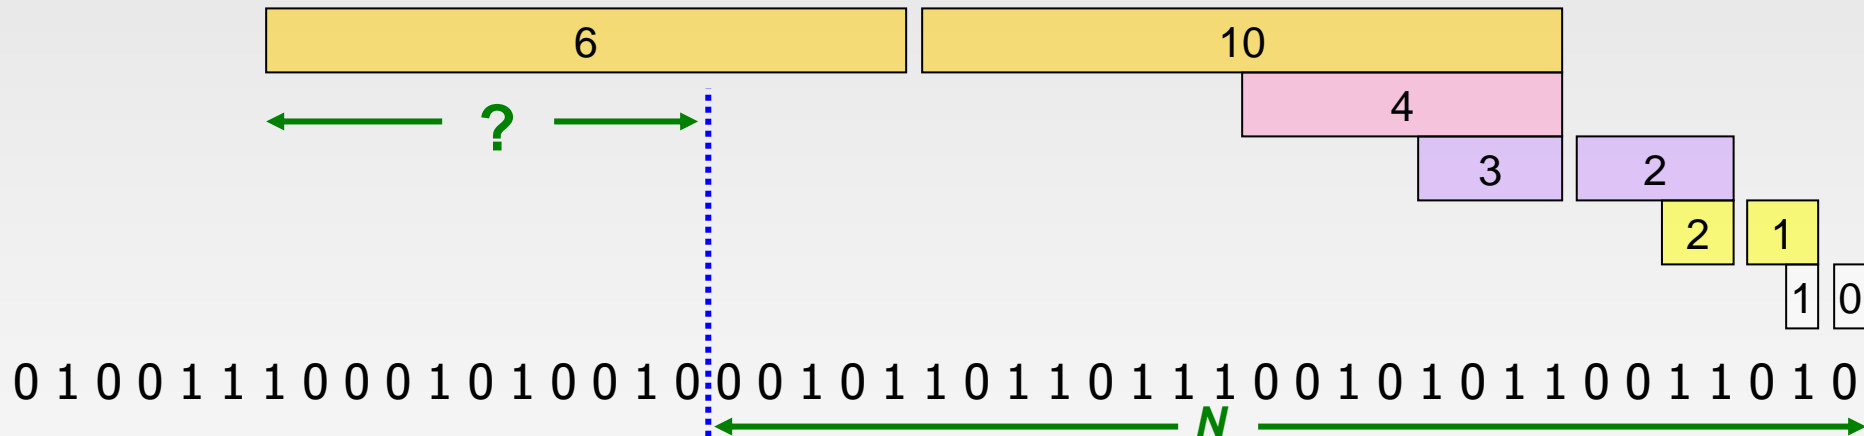
We can reconstruct the count of the last N bits, except we are not sure how many of the last 6 1s are included in the N

What's Good?

- ❖ Stores only $O(\log^2 N)$ bits
 - $O(\log N)$ counts of $\log_2 N$ bits each
- ❖ Easy update as more bits enter
- ❖ Error in count no greater than the number of **1s** in the “**unknown**” area


What's Not So Good?

- ❖ As long as the 1s are fairly evenly distributed, the error due to the unknown region is small – **no more than 50%**
- ❖ But it could be that all the 1s are in the unknown area at the end
- ❖ In that case, **the error is unbounded!**
 - Because that the number of 1's in the known regions could be 0!



Fixup: DGIM Algorithm

- ❖ **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
 - Let the block **sizes** (number of **1s**) increase exponentially
- ❖ When there are few 1s in the window, block sizes stay small, so errors are small




DGIM: Timestamps

- ❖ Each bit in the stream has a timestamp, starting from 1, 2, ...
- ❖ Record timestamps modulo N (the window size), so we can represent any relevant timestamp in $O(\log_2 N)$ bits
 - E.g., given the windows size 40 (N), timestamp 123 will be recorded as 3, and thus the encoding is on 3 rather than 123

DGIM: Buckets


- ❖ A bucket in the DGIM method is a record consisting of:
 - (A) The timestamp of its end [$O(\log N)$ bits]
 - (B) The number of 1s between its beginning and end [$O(\log \log N)$ bits]
- ❖ Constraint on buckets:
 - Number of 1s must be a power of 2
 - That explains the $O(\log \log N)$ in (B) above



Representing a Stream by Buckets

- ❖ The right end of a bucket is always a position with a 1
- ❖ Every position with a 1 is in some bucket
- ❖ Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- ❖ Buckets do not overlap in timestamps
- ❖ **Buckets are sorted by size**
 - Earlier buckets are not smaller than later buckets
- ❖ Buckets disappear when their end-time is $> N$ time units in the past

Example: Bucketized Stream



- ❖ Three properties of buckets that are maintained:
 - Either **one** or **two** buckets with the same power-of-2 number of 1s
 - Buckets do not overlap in timestamps
 - Buckets are sorted by size

Updating Buckets

- ❖ When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to ***N*** time units before the current time
- ❖ 2 cases: Current bit is **0** or **1**
- ❖ If the current bit is 0: no other changes are needed
- ❖ If the current bit is 1:
 - (1) Create a new bucket of size 1, for just this bit
 - ▶ End timestamp = current time
 - (2) If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
 - (3) If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
 - (4) And so on ...

Example: Updating Buckets

Current state of the stream:

10010101100010110 101010101010110 1010101010101110 1010101110101000 10110010

Bit of value 1 arrives

0010101100010110 101010101010110 1010101010101110 1010101110101000 101100101

Two white buckets get merged into a yellow bucket

0010101100010110 101010101010110 1010101010101110 1010101110101000 101100101

Next bit 1 arrives, new orange white is created, then 0 comes, then 1:

0101100010110 101010101010110 1010101010101110 1010101110101000 101100101101

Buckets get merged...

0101100010110 101010101010110 1010101010101110 1010101110101000 101100101101

State of the buckets after merging

0101100010110 10101010101011010101010101110 1010101110101000 101100101101

How to Query?

- ❖ To estimate the number of 1s in the most recent N bits:
 - Sum the sizes of all buckets but the last
 - ▶ (note “size” means the number of 1s in the bucket)
 - Add half the size of the last bucket
- ❖ Remember: We do not know how many 1s of the last bucket are still within the wanted window
- ❖ Example:


At least 1 of
size 16. Partially
beyond window.

2 of
size 8

2 of
size 4


1 of
size 2

2 of
size 1



Error Bound: Proof

- ❖ Why is error 50%? Let's prove it!
- ❖ Suppose the last bucket has size 2^r
- ❖ Then by assuming 2^{r-1} (i.e., half) of its 1s are still within the window, we make an error of at most 2^{r-1}
- ❖ Since there is at least one bucket of each of the sizes less than 2^r , the true sum is at least
 $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- ❖ Thus, error at most 50%



Further Reducing the Error

- ❖ Instead of maintaining 1 or 2 of each size bucket, we allow either $r-1$ or r buckets ($r > 2$)
 - Except for the largest size buckets; we can have any number between 1 and r of those
- ❖ Error is at most $O(1/r)$
 - *WHY?*
- ❖ By picking r appropriately, we can tradeoff between number of bits we store and the error

Practice


- ❖ Suppose we are maintaining a count of 1s using the DGIM method. We represent a bucket by (i, t) , where i is the number of 1s in the bucket and t is the bucket timestamp (time of the most recent 1).
 - Consider that the current time is 200, window size is 60, and the current list of buckets is: $(16, 148)$ $(8, 162)$ $(8, 177)$ $(4, 183)$ $(2, 192)$ $(1, 197)$ $(1, 200)$. At the next ten clocks, 201 through 210, the stream has 0101010101. What will the sequence of buckets be at the end of these ten inputs?

Solution

- ❖ There are 5 1s in the stream. Each one will update to windows to be:
 - (1) (16, 148)(8, 162)(8, 177)(4, 183)(2, 192)(1, 197)(1, 200), (1, 202)
=> (16, 148)(8, 162)(8, 177)(4, 183)(2, 192)(2, 200), (1, 202)
 - (2) (16, 148)(8, 162)(8, 177)(4, 183)(2, 192)(2, 200), (1, 202), (1, 204)
 - (3) (16, 148)(8, 162)(8, 177)(4, 183)(2, 192)(2, 200), (1, 202), (1, 204), (1, 206)
=> (16, 148)(8, 162)(8, 177)(4, 183)(2, 192)(2, 200), (2, 204), (1, 206)
=> (16, 148)(8, 162)(8, 177)(4, 183)(4, 200), (2, 204), (1, 206)
 - (4) Windows Size is 60, so (16,148) should be dropped.
(16, 148)(8, 162)(8, 177)(4, 183)(4, 200), (2, 204), (1, 206), (1, 208) =>
(8, 162)(8, 177)(4, 183)(4, 200), (2, 204), (1, 206), (1, 208)
 - (5) (8, 162)(8, 177)(4, 183)(4, 200), (2, 204), (1, 206), (1, 208), (1, 210)
=> (8, 162)(8, 177)(4, 183)(4, 200), (2, 204), (2, 208), (1, 210)

Extensions

- ❖ Answer queries **How many 1's in the last k ?** where $k < N$
 - **A:** Find earliest bucket **B** that overlaps with k .
Number of **1s** is the **sum of sizes of more recent buckets + $\frac{1}{2}$ size of B**



- ❖ Can we handle the case where the stream is not bits, but integers, and we want the sum of the last k elements?


Extensions

- ❖ Stream of positive integers
- ❖ We want the sum of the last k elements

➤ Amazon: Avg. price of last k sales

❖ Solution:

- (1) If you know all have at most m bits
 - ▶ Treat m bits of each integer as a separate stream
 - ▶ Use DGIM to count 1s in each integer
 - ▶ The sum is $= \sum_{i=0}^{m-1} c_i 2^i$ c_i ...estimated count for the i -th bit
- (2) Use buckets to keep partial sums
 - ▶ Sum of elements in size b bucket is at most 2^b



Idea: Sum in each bucket is at most 2^b (unless bucket has only 1 integer)

Bucket sizes:



References

- ❖ Chapter 4, Mining of Massive Datasets.

End of Chapter 6.1