

# 9b. Vision and Language Reasoning

Never Stand Still

Faculty of Engineering

COMP9444 Week 9b

Sonit Singh

School of Computer Science and Engineering

Faculty of Engineering

The University of New South Wales, Sydney, Australia

[sonit.singh@unsw.edu.au](mailto:sonit.singh@unsw.edu.au)

# **WARNING**

This material has been reproduced and communicated to you  
by or on behalf of the University of New South Wales in  
accordance with section 113P(1) of the Copyright Act 1968 (Act).

The material in this communication may be subject to  
copyright under the Act. Any further reproduction or  
communication of this material by you may be the subject of  
copyright protection under the Act.

Do not remove this notice

# Goal


- To motivate the need for “Computer Vision + Natural Language Processing”
- After the session, everyone can confidently say: “yeah, I know various tasks at the intersection of computer vision and natural language processing”
- Focus on high-level overview, not technical details
- Focus on static images, not videos (although they are easy to translate to videos)
- Focus on selective set of papers for various tasks, not a comprehensive literature review

# Agenda

- Computer Vision
- Natural Language Processing
- Computer Vision + Natural Language Processing
- Building Blocks
  - Convolutional Neural Networks (CNNs)
  - Recurrent Neural Networks (RNNs)
  - Attention Mechanism
- Encoder-Decoder Framework
- Image Captioning
- Visual Question Answering (VQA)
- Visual Dialog (VisDial)
- Vision-Language Navigation (VLN)
- Visual Grounding
- Summary


# Computer Vision

- Enabling machines to process, represent, understand, and generate visual data



# CV Applications

## ➤ Image Classification




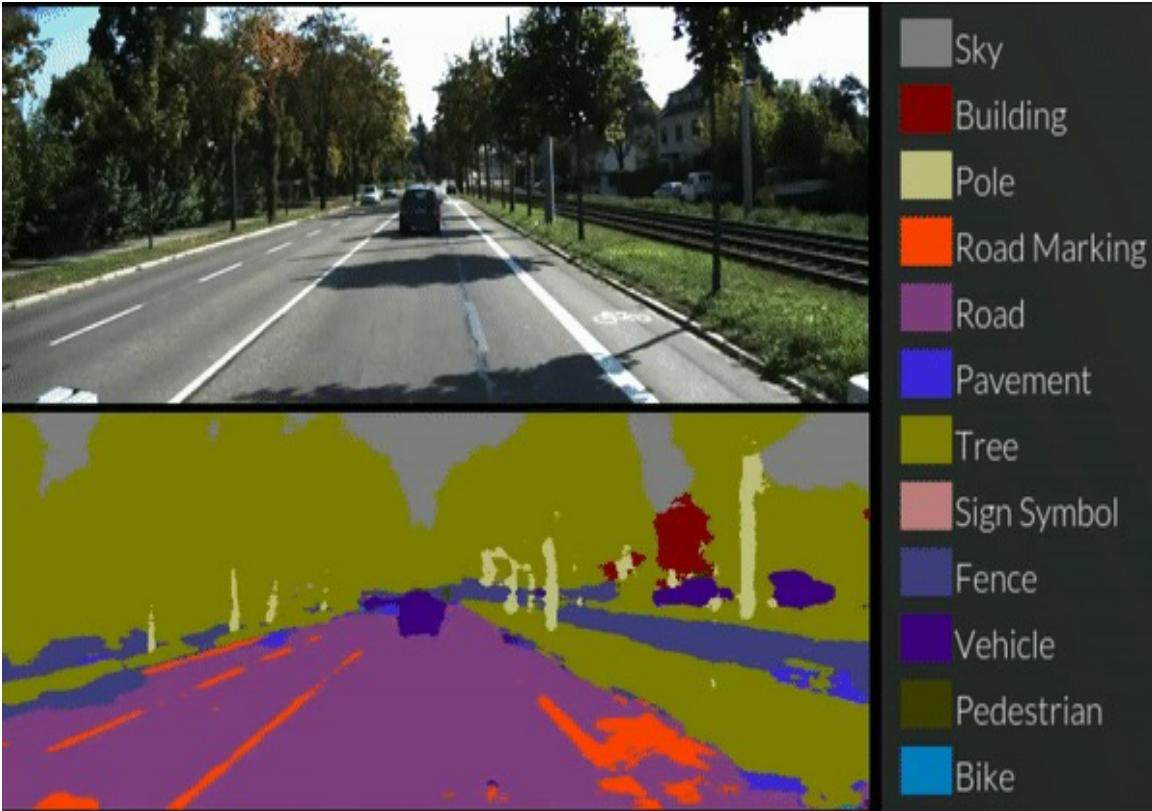
# CV Applications

## ➤ Object Detection



# CV Applications

## ➤ Segmentation



# Natural Language Processing

- Enabling machines to process, represent, understand, and generate languages

In fact, the Chinese NORP market has the three CARDINAL most influential names of the retail and tech space – Alibaba GPE , Baidu ORG , and Tencent PERSON (collectively touted as BAT ORG ), and is betting big in the global AI GPE in retail industry space . The three CARDINAL giants which are claimed to have a cut-throat competition with the U.S. GPE (in terms of resources and capital) are positioning themselves to become the ‘future AI PERSON platforms’. The trio is also expanding in other Asian NORP countries and investing heavily in the U.S. GPE based AI GPE startups to leverage the power of AI GPE . Backed by such powerful initiatives and presence of these conglomerates, the market in APAC AI is forecast to be the fastest-growing one CARDINAL , with an anticipated CAGR PERSON of 45% PERCENT over 2018 - 2024 DATE .

To further elaborate on the geographical trends, North America LOC has procured more than 50% PERCENT of the global share in 2017 DATE and has been leading the regional landscape of AI GPE in the retail market. The U.S. GPE has a significant credit in the regional trends with over 65% PERCENT of investments (including M&As, private equity, and venture capital) in artificial intelligence technology. Additionally, the region is a huge hub for startups in tandem with the presence of tech titans, such as Google ORG , IBM ORG , and Microsoft ORG .

# NLP Applications

## ➤ Text Understanding

"I love this movie.  
I've seen it many times  
and it's still awesome."




"This movie is bad.  
I don't like it at all.  
It's terrible."



# NLP Applications

## ➤ Machine Translation



The screenshot shows the Google Translate interface. At the top, there are tabs for "Text" and "Documents". Below that, the source language is set to "ENGLISH" and the target language is "GERMAN". The input text on the left is "I love teaching humans and machines". The output text on the right is "Ich liebe es, Menschen und Maschinen beizubringen". There are icons for microphone, speaker, and sharing below the text boxes. A progress bar at the bottom indicates 35 / 5000. On the far right, there are "Send feedback" and other sharing options.

# NLP Applications

- Question Answering/Comprehension

## Passage Sentence

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under gravity.

## Question


What causes precipitation to fall?

## Answer Candidate

gravity


# NLP Applications

## ➤ Text Summarization



# Vision + Language

- Science Perspective
  - Vision is how we observe and understand the world
  - Language is how we communicate
- A move towards Artificial General Intelligence (AGI)




# Applications at the intersection of Vision and Language

- To aid “visually impaired” people



# Applications at the intersection of Vision and Language

- To aid “situationally impaired” analysts



Did anyone enter this room last week?



Yes, 127 instances logged on camera

Show me images of anyone carrying a black bag.



...

# Applications at the intersection of Vision and Language

## ➤ Personal Assistants



Add the chopped zucchini to the pan...

And what's next?

Hey, can you order more zucchini?

Done.



# Applications at the intersection of Vision and Language

- Natural Language Instructions for Robots



Is there smoke in any room around you?

Yes, in one room




Go there and look for people

...


# Applications at the intersection of Vision and Language

- Generating and summarizing radiology reports from medical images



# Building Blocks: Convolutional Neural Networks (CNNs)


- A class of deep neural networks suitable for processing 2D/3D data. For e.g., Images and Videos
- CNNs can capture high-level representation of images/videos which can be used for end-tasks such as classification, object detection, segmentation, etc.
- A range of CNNs improving over the years



# CNN Architecture

- A typical CNN architecture consists of the following layers:


- Convolution layer
- ReLU layer (non-linearity)
- Pooling layer
- Flattening
- Fully-connected layer
- Output layer



- There can be multiple steps of convolution followed by pooling, before reaching the fully connected layers.

# Building Blocks: Recurrent Neural Networks (RNNs)

- A class of neural networks suitable for processing temporal or sequential data
- The basic unit of RNN is called “cell”, and each cell consists of layers and a series of cells that enables the sequential processing of recurrent neural network models
- RNNs have a looping mechanism that acts as a highway to allow information to flow from one step to the next. This information is the hidden state, which is a representation of previous inputs.




# Long-range Dependency problem

- Vanilla RNNs suffers from vanishing gradient problem
  - As the RNNs processes more steps, it has troubles retaining information from previous steps.
  - Due to back-propagation, the earlier layers fail to do any learning as the internal weights are barely being adjusted due to extremely small gradients.
  - Does not learn the long-range dependencies across time steps



“Once upon a time, there was a king who ruled a great and glorious nation. Favourite amongst his subjects was the court painter of whom he was very proud. Everybody agreed this wizened old man painted the greatest pictures in the whole kingdom and the king would spend hours each day gazing at them in wonder. However, one day a dirty and disheveled stranger presented himself at the court claiming that in fact he was the greatest painter in the land. The indignant king decreed a competition would be held between the two artists, confident it would teach the vagabond an embarrassing lesson. Within a month they were both to produce a masterpiece that would out do the other. After thirty days of working feverishly day and night, both artists were ready. They placed their paintings, each hidden by a cloth, on easels in the great hall of the castle. As a large crowd gathered, the king ordered the cloth be pulled...”

# LSTMs/GRUs

- LSTMs and GRUs are two special RNNs, capable of learning long-term dependencies using mechanisms called **gates**.
- These gates are different tensor operations that can learn what information to add or remove to the hidden state.




# End-to-End text classification




# Seq2Seq model

- Model has two parts: Encoder and Decoder (Encoder-Decoder Framework)




# Seq2Seq used for various NLP applications

- Machine Translation




- Automatic Email Reply





# Issues with RNNs for seq2seq tasks

- Despite being very successful for various NLP tasks, there were challenges:
  - dealing with long-range dependencies
  - the sequential nature of the architecture prevents parallelization
- Attention mechanism helped to overcome first issue to certain extent



# Attention Mechanism

- A set of mechanisms that limit some processing to a subset of incoming stimuli (reducing computational demands)
- Attention in neural networks
  - A mechanism that **learns to focus** on a subset of the **input** that is **relevant to the task**.





# Neural Machine Translation by jointly Learning to Align and Translate

- Issue: The encoder-decoder framework compresses all the necessary information of a source sentence into a fixed-length vector.
- Solution: Enable the network to pay attention to specific areas of the input by adding new (weighted) connections

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$


$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})},$$

$$e_{ij} = a(s_{i-1}, h_j)$$




# NMT by jointly Learning to Align and Translate

**Before Attention:** Long sentences are very hard as they are compressed to a fixed length vector



**After Attention:** The attention mechanism helps to overcome the issue



# Image Captioning

- Developing models to generate textual description of an image



"man in black shirt is playing guitar."





"construction worker in orange safety vest is working on road."




"two young girls are playing with lego toy."

# Show and Tell (Vinyals et al., 2015)



# Show, Attend and Tell (Xu et al., 2015)

- “Rather than compressing an entire image into static representation, attention allows for salient features to dynamically come to forefront as needed”



# Show, Attend and Tell (Xu et al., 2015)



A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.



<start>



a



large



airplane



flying



in



the



blue



sky




<end>




# Image Captioning with Semantic Attention (You et al., 2016)

- Learns to selectively attend to semantic concept proposals and fuse them into hidden states and outputs of RNNs.




# Image Captioning with Semantic Attention (You et al., 2016)

- Learns to selectively attend to semantic concept proposals and fuse them into hidden states and outputs of RNNs.

|   |   |  |   |   |   |   |   |
|---|---|--|---|---|---|---|---|
|  |  |  |  |  |  |  |  |
| Google NIC  | a white plate topped with a variety of food.                                      | a baby is eating a piece of paper.   | a close up of a plate of food on a table.   | a teddy bear sitting on top of a chair .  | a person is holding colorful umbrella.  | a woman is holding a cell phone in her hand .                                       | a traffic light is on a city street.  |
| Top-5 visual attributes   | plate broccoli fries food french  | teeth brushing toothbrush holding baby   | cake table plate sitting birthday   | teddy cat bear stuffed white  | umbrella beach water sitting boat   | woman bathroom her scissors man   | street sign cars clock traffic  |
| ATT-FCN   | a plate with a sandwich and french fries.   | a baby with a toothbrush in its mouth.   | a table topped with a cake with candles on it.                                      | a white teddy bear sitting next to a stuffed animal .                               | a black umbrella sitting on top of a sandy beach .                                  | a woman holding a pair of scissors in her hands .                                   | a street with cars and a clock tower next to a building.                            |

# Knowing When to Look (Lu et al., 2017)

- Words such as “a”, “of”, “it” may be seen as not worth attending
- Words such as “woman”, “dog”, “traffic light” need attending to the image
- Automatically determines when to look (sentinel gate) and where to look (spatial attention) for word generation.






# Knowing When to Look (Lu et al., 2017)

- Visualization of generated captions and image attention maps on COCO dataset



# Bottom-up and Top-down attention (Anderson et al., 2018)


- Previous attention models operate on CNN features corresponding to a uniform grid of equally-sized image regions
- Bottom-up and Top-down enables attention to be calculated at the level of objects (or salient image regions)
- Features vectors extracted from Faster R-CNN are used



# Bottom-up and Top-down attention (Anderson et al., 2018)



Ours: Resnet – A man sitting on a *toilet* in a bathroom.



Ours: Up-Down – A man sitting on a *couch* in a bathroom.




## And many more work ...

- Show, Control and Tell: A Framework for Generating Controllable and Grounded Captions – Cornia M et al., CVPR 2019
- Pointing Novel Objects in Image Captioning, Li Y et al., CVPR 2019
- Describing like humans: on diversity in image captioning – Wang Q et al., CVPR 2019
- Length-Controllable Image Captioning – Deng C et al., ECCV 2020
- Transformer-based local-global guidance for image captioning – Parvin et al., 2023
- Show, tell and summarise: learning to generate and summarise radiology findings from medical images – Singh et al., Neural Computing & Applications, 2021.
- Medical image captioning via generative pretrained transformers – Selivanov et al. Scientific Reports, 2023
- Let's see Image captioning in action
  - <https://milhidaka.github.io/chainer-image-caption/>

# Visual Question Answering (VQA)

- Given an **image**, can our machine answer the corresponding **questions in natural language**?



# VQA Dataset



Who is wearing glasses?

man



woman




Where is the child sitting?

fridge



arms



Is the umbrella upside down?

yes



no



How many children are in the bed?

2




1




# VQA Approach: Bag-of-words + Image feature (iBOWIMG)

- Combine image and word embeddings to predict answer




# VQA Common Approach


- Combine CNN (for vision) and RNN (for language) to predict answer




# VQA – set of problems




(a) "Yes/no" problem



(b) Multi-choice problem



(c) Number counting problem




(d) Open-ended problem

Figure 2 Common types of visual question answering. "Yes/No" problem and multi-choice problem can be regarded as a classification problem, while number counting problem and open-ended problem can be viewed as a caption generation problem.

# VQA in the Medical Domain


- Aims to create a system that can answer natural language questions based on a given medical image.
- VQA-Med 2019 Challenge dataset




Q: What is the plane of this MRI?  
A: Sagittal



Q: The CT scan shows what organ system?  
A: Spine and contents




Q: With what modality is this image taken?  
A: AN - angiogram



Q: What is most alarming about this MRI?  
A: Schwannoma


# VQA-Med

- Most questions in VQA-Med 2019 dataset are “close-ended”
- More than 50% of answers consists of only one word, and more than 82% of answers have between one and three words.
- Best strategy is to do classification rather than generation



# VQA-Med Methodology

- Contributions
  - Incorporating medical domain knowledge
    - Image encoder – applied self-supervised pretraining using Radiology Objects in COntext (ROCO) dataset
    - Question encoder – used BioBERT, pretrained on same tasks as BERT, but using the PubMed corpus
- Evidence verification
  - Gradient Weighted Class Activation Map (Grad-CAM) for evidence verification



# Results

- Test accuracy achieved by each model variation


| Model Variation                                 | Test Accuracy |
|---|---------------|
| VGG-16 + LSTM + Concatenation                   | 0.56          |
| ResNet-50 + LSTM + Concatenation                | 0.54          |
| ResNet-152 + LSTM + Concatenation               | 0.53          |
| <b>VGG-16 + BERT + Concatenation</b>            | <b>0.60</b>   |
| VGG-16 + BERT + SAN                             | 0.58          |
| <b>VGG-16 + BioBERT + Concatenation</b>         | <b>0.60</b>   |
| <b>Pretrained VGG-16 + BERT + Concatenation</b> | <b>0.60</b>   |

- Accuracy of the baseline vs. BERT model per category type

| Model Variation | Baseline | +BERT       |
|-----------------|----------|-------------|
| Modality        | 0.64     | 0.76        |
| Plane           | 0.78     | 0.77        |
| Organ           | 0.74     | 0.74        |
| Abnormality     | 0.06     | 0.08        |
| Overall         | 0.56     | <b>0.60</b> |

# Results

## ➤ Effect of using BERT vs. LSTM



Q: What imaging modality was used to take this image?  
GT: CT with IV contrast  
Baseline: **Skull fracture from cell phone**  
+BERT: CT with IV contrast

(a) Category misclassification



Q: What was this image taken with?  
GT: MR – PDW proton density  
Baseline: Yes  
+BERT: MR – PDW proton density

(b) Wrong answer type



Q: Is this a contrast or noncontrast MRI?  
GT: Noncontrast  
Baseline: **MR – flair**  
+BERT: Noncontrast

(c) Not choosing from options


# Evidence verification

- Attention distribution output by Stacked Attention Network (SAN) fusion method




Q: What imaging modality is seen here?  
GT: XR – plain film  
A: XR – plain film

(a)



Q: What plane is demonstrated?  
GT: Axial  
A: Coronal


(b)



Q: What is most alarming about this MRI?  
GT: Hypothalamic hamartoma  
A: Meningioma


(c)

- Grad-CAM output




Q: Is this a T1 weighted image?  
GT: No  
A: No

(a)




Q: In what plane is this MRI captured?  
GT: Sagittal  
A: Sagittal

(b)




Q: What part of the body is being imaged?  
GT: Gastrointestinal  
A: Gastrointestinal

(c)



Q: What is abnormal in the X-ray?  
GT: Chondroblastoma  
A: Jones fracture and dancer's fracture

(d)



Q: What abnormality is seen in the image?  
GT: Spondylolisthesis bilateral pars fracture  
A: Spondylolysis


(e)

# Let's see in action

- <https://visualqa.org/visualize/>
- <https://demo.allennlp.org/visual-question-answering>

# Visual Dialog

- Ability to hold a meaningful dialog with humans in natural language about visual content
- Given an image  $I$ , a history of dialog consisting of sequence of question-answer pairs, and a natural language **follow-up question**, the task for the machine is to answer the question in free-form natural language.



# Image Captioning vs. VQA vs. Visual Dialog



## VQA

Q: How many people on wheelchairs ?

A: Two

Q: How many wheelchairs ?

A: One

## Captioning

Two people are in a wheelchair and one is holding a racket.



## Visual Dialog

Q: How many people are on wheelchairs ?

A: Two

Q: What are their genders ?

A: One male and one female

Q: Which one is holding a racket ?

A: The woman

Session Variables

Q: What is the gender of the one in the white shirt ?

A: She is a woman

Q: What is she doing ?

A: Playing a Wii game

Q: Is that a man to her right

A: No, it's a woman

# Visual Dialog – Late Fusion Encoder

- Entire history  $H$  is concatenated and encoded by LSTM




Image I

Do you think the woman is with him?

Question  $Q_t$


The man is riding his bicycle on the sidewalk. Is the man wearing a helmet? No he does not have a helmet on. ... Are there any people nearby? Yes there's a woman walking behind him.

$t$  rounds of history  
(concatenated)




# Visual Dialog – Hierarchical Recurrent Encoder

- Each QA-pairs in history  $H$  separately encoded by the LSTM



# Visual Dialog – Hierarchical Recurrent Encoder

- Compute inner product of question vector with each history vector to get scores over previous rounds, which are fed to a SoftMax to get attention-over-history probabilities.



# Improving Visual Dialog (and few examples)

➤ Generates more diverse, image relevant and fluent dialog.

Image + Caption



a cutting board with carrots next to a cell phone.



many brown and black cows standing in the grass



many elephants standing in an enclosure near a building



a reflection of the tub is in the bathroom mirror along with the sink , you can see all the room

GT dialog (Das et al., 2017a)

|                                      |   |
|--------------------------------------|---|
| what color is the cutting board?     | blue                                      |
| how big is it                        | about the size of a regular laptop screen |
| how many carrots are there ?         | 4   |
| are they the orange kind ?           | 2 are orange , 2 are white                |
| what kind of cell phone is it ?      | a samsung                                 |
| can you see any text on the screen ? | a recipe for UNK UNK soup                 |
| do you see the room ?                | just the counter                          |
| are there people ?                   | not visible                               |
| what color is the counter ?          | white                                     |

Q-BOT-A-BOT dialog

|                              |           |
|------------------------------|-----------|
| is the photo in color ?      | no        |
| how many people do you see ? | no people |
| is this a close up ?         | no        |
| is this a close up ?         | zebra is  |
| is this a close up ?         | no        |
| is this a close up ?         | no        |
| is this a close up ?         | no        |
| is this a close up ?         | no        |
| is this a close up ?         | no        |
| is this a close up ?         | no        |
| is this a close up ?         | no        |

Diverse-Q-BOT-A-BOT dialog

|  |                                 |
|--|---------------------------------|
| is the photo in color ?                  | 1 picture is in color           |
| is the photo close up ?                  | black and white                 |
| is the photo in color ?                  | no it 's black and white        |
| is it a close up ?                       | no it 's a close up of the cake |
| can you tell what brand it is ?          | no i ca n't see the brand       |
| can you tell if it 's a male or female ? | no i ca n't tell                |
| can you tell if it 's a male or female ? | no i ca n't tell                |
| is this in a kitchen ?                   | nope it 's in a home            |
| are there any utensils ?                 | no                              |
| are there any utensils ?                 | nope                            |

|                               |                    |
|-------------------------------|--------------------|
| do you see any people ?       | 1                  |
| do you see a fence ?          | a small 1          |
| how many cows are there ?     | 5                  |
| what color are the cows ?     | black and white    |
| can you see the sky ?         | some               |
| can you see the sky ?         | nope               |
| are there any other animals ? | not that i can see |
| can you see the sky ?         | not at all         |
| are there any trees ?         | in the background  |
| can you see the sky ?         | not at all         |

|                                |                                   |
|--------------------------------|-----------------------------------|
| how many elephants are there ? | 1                                 |
| any people ?                   | ca n't see any                    |
| male or female ?               | not sure                          |
| what is the elephant doing ?   | just standing there               |
| male or female ?               | there are 2 males and 1 female    |
| can you see the sky ?          | 1 of it                           |
| what color is the building ?   | just a black and white picture    |
| do you see any trees ?         | there are trees in the background |
| are there any other animals ?  | red and white                     |
| do you see any trees ?         | i see some in the background      |

|                             |                        |
|-----------------------------|------------------------|
| is the photo in color ?     | no                     |
| is there any people ?       | 5                      |
| what color walls ?          | i ca n't see the walls |
| is there a mirror ?         | yes                    |
| is there a mirror ?         | yes                    |
| is there a mirror ?         | yes                    |
| are the walls seen ?        | no                     |
| what color are the walls ?  | no walls               |
| what is the floor made of ? | yes                    |
| what color is the floor ?   | i ca n't see the floor |

# Many more tasks at CV+NLP

## ➤ Vision-Language Navigation (VLN)

### Instruction

Turn right and head towards the *kitchen*. Then turn left, pass a *table* and enter the *hallway*. Walk down the hallway and turn into the *entry way* to your right *without doors*. Stop in front of the *toilet*.



Initial Position



Target Position



Demonstration Path A



Executed Path B



Executed Path C





Figure 1: Demonstration of the VLN task. The instruction, the local visual scene, and the global trajectories in a top-down view. The agent does not have access to the top-down view. Path A is the demonstration path following the instruction. Path B and C represent two different paths executed by the agent. Figure credit: Wang et al. (2019).

# Many more tasks at CV+NLP


- Visual grounding for referring expressions
- Grounding an utterance to refer to something or someone in image



(a)



(b)



(c)

Fig. 1: Interactive visual grounding of referring expressions. (a) Ground self-referential expressions. (b) Ground relational expressions. (c) Ask questions to resolve ambiguity. Red boxes indicate referred objects. Blue dashed boxes indicate candidate objects. See also the accompanying video at <http://bit.ly/INGRESSvid>.

- Let's see video in action: <http://bit.ly/INGRESSvid>

# Many more tasks at CV+NLP

- Fun stuff: Comic books





Figure 5. The image-text architecture applied to an instance of the *text cloze* task. Pretrained image features are combined with learned text features in a hierarchical LSTM architecture to form a context representation, which is then used to score text candidates.

# Many more tasks at CV+NLP

## ➤ From Recognition to Cognition: Visual Commonsense Reasoning



hide all

show all

[person1]

[person2]

[person3]

[person4]

more objects »

Why is [person4] pointing at [person1]?

- a) He is telling [person3] that [person1] ordered the pancakes.
- b) He just told a joke.
- c) He is feeling accusatory towards [person1].
- d) He is giving [person1] directions.

Rationale: I think so because...

- a) [person1] has the pancakes in front of him.
- b) [person4] is taking everyone's order and asked for clarification.
- c) [person3] is looking at the pancakes both she and [person2] are smiling slightly.
- d) [person3] is delivering food to the table, and she might not know whose order is whose.

# Summary

- Vision + Language: Help us to understand human brain functioning better
- Visual Turing Test for modern AI systems
- A step towards Artificial General Intelligence (AGI)
- Variety of applications:
  - Helping visually-impaired people
  - Early child education
  - Personal assistants
  - Robot navigation
  - Video surveillance systems
  - Search engines




**Tell us about your  
experience. Shape  
the future of  
education at  
UNSW.**

Visit [Moodle](#) to complete  
the myExperience survey



**UNSW**  
SYDNEY



**Tell us about your  
experience.  
Shape the future of  
education at UNSW.**

Visit [Moodle](#) to complete  
the myExperience survey

Please be mindful of the [UNSW Student Code of Conduct](#) as you provide feedback. At UNSW we aim to provide a respectful community and ask you to be careful to avoid any language that is sexist, racist or likely to be hurtful. You should feel confident that you can provide both positive and negative feedback, but please be considerate in how you communicate.



**UNSW**  
SYDNEY



**UNSW**  
SYDNEY



## Questions?

If you feel fascinated and want to get involved in research, feel free to reach out at  
<https://research.unsw.edu.au/people/dr-sonit-singh>